

PA01: Processes and Threads

Graduate Systems (CSE638)

Submitted by:

Adarsh Singh

Roll Number: PhD25001

IIIT Delhi

January 23, 2026

Table of Contents

1. Introduction
2. Part A: Program Implementation
3. Part B: Worker Functions
4. Part C: Measurement Script and Results
5. Part D: Scaling Analysis
6. Observations and Analysis
7. AI Usage Declaration
8. Conclusion
9. Appendix: Complete Source Code

1. Introduction

This report documents the implementation and analysis of PA01 assignment for Graduate Systems (CSE638). The assignment explores the differences between process-based and thread-based parallelism through practical implementation and measurement of CPU-intensive, memory-intensive, and I/O-intensive workloads.

1.1 Objectives

The primary objectives of this assignment are:

1. Implement Program A using fork() to create child processes
2. Implement Program B using pthread to create threads
3. Create three worker functions: cpu, mem, and io
4. Measure and compare CPU%, Memory%, and I/O metrics
5. Analyze scaling behavior as process/thread count increases

1.2 System Configuration

The experiments were conducted on a Linux system with the following tools:

- top: For CPU and memory utilization monitoring
- taskset: For CPU affinity/pinning
- iostat: For disk I/O statistics
- /usr/bin/time: For execution time measurement

2. Part A: Program Implementation

2.1 Program A: Process-based (fork)

Program A creates child processes using the fork() system call. Each child process executes the specified worker function independently with its own address space.

Line-by-Line Explanation

Lines 1-16: Header and Includes

```
#include <stdio.h>      // Standard I/O functions
#include <stdlib.h>     // exit(), atoi(), malloc()
#include <unistd.h>     // fork(), getpid()
#include <sys/wait.h>   // wait()
#include <math.h>        // sin(), cos()
#include <fcntl.h>       // open(), O_WRONLY, O_CREAT
```

These headers provide essential system calls for process creation, file operations, and mathematical functions.

Lines 18-23: Constants Definition

```
#define BASE_LOOP_COUNT 1000          // 1 (last digit) × 103
#define CPU_LOOP_COUNT 1000000000    // 100M for visible CPU burst
#define MEM_ARRAY_SIZE 10000000     // ~40MB per process
#define IO_FILE_SIZE 1024           // 1KB buffer
```

BASE_LOOP_COUNT is derived from roll number PhD25001 (last digit = 1, so $1 \times 1000 = 1000$). CPU_LOOP_COUNT is increased to 100 million to ensure the CPU burst lasts long enough for 'top' to capture it.

Lines 75-97: Main Function

```
int num_processes = (argc >= 3) ? atoi(argv[2]) : 2;
```

Default is 2 processes as required by the assignment.

```
for (int i = 0; i < num_processes; i++) {
    pid_t pid = fork();
    if (pid == 0) { /* child executes worker */ exit(0); }
}
```

The fork() call creates a child process. In the child (pid == 0), the appropriate worker function is executed, then the child exits. The parent continues the loop to create more children.

```
for (int i = 0; i < num_processes; i++) wait(NULL);
```

Parent waits for all children to complete using wait().

2.2 Program B: Thread-based (pthread)

Program B creates threads using the POSIX pthread library. Threads share the same address space, making communication easier but requiring careful synchronization.

Line-by-Line Explanation

Lines 20-23: Thread Argument Structure

```
typedef struct {  
    int id;  
    const char *type;  
} thread_arg_t;
```

A structure to pass arguments to each thread, containing the thread ID and worker type.

Lines 81-87: Thread Creation and Joining

```
pthread_create(&threads[i], NULL, worker_func, &t_args[i]);
```

Creates a new thread that executes worker_func with the given arguments.

```
pthread_join(threads[i], NULL);
```

Waits for each thread to complete before the program exits.

3. Part B: Worker Functions

3.1 CPU Worker Function

Definition: CPU-intensive programs spend the majority of their execution time performing calculations on the CPU, rather than waiting for data from other resources.

Implementation:

```
void cpu_worker(void) {
    volatile double result = 0.0;
    for (long i = 0; i < CPU_LOOP_COUNT; i++) {
        result += sin((double)i) * cos((double)i);
        if (i % 1000 == 0) result *= 1.000001;
    }
}
```

Line-by-Line Explanation:

- volatile double result: Prevents compiler optimization that might skip calculations
- sin() * cos(): Complex floating-point operations that stress the CPU
- 100 million iterations: Ensures the burst is long enough to measure
- result *= 1.000001: Additional operation to prevent loop optimization

3.2 Memory Worker Function

Definition: Memory-intensive programs are bottlenecked by the speed and capacity of the system's memory (RAM). They require large amounts of data to be moved between CPU and memory.

Implementation:

```
void mem_worker(void) {
    int *large_array = malloc(MEM_ARRAY_SIZE * sizeof(int));
    for (int i = 0; i < MEM_ARRAY_SIZE; i++) large_array[i] = i;
    for (int iter = 0; iter < BASE_LOOP_COUNT * 50; iter++) {
        int idx = rand_r(&seed) % MEM_ARRAY_SIZE;
        sum += large_array[idx];
        large_array[idx] = sum;
    }
    free(large_array);
}
```

Line-by-Line Explanation:

- malloc(40MB): Allocates a large array (~40MB) to stress memory
- Random access pattern: rand_r() generates random indices, defeating CPU cache

- Read and write operations: Both reading ($\text{sum} += \text{arr}[\text{idx}]$) and writing ($\text{arr}[\text{idx}] = \text{sum}$) stress memory
- `rand_r()`: Thread-safe random number generator

3.3 I/O Worker Function

Definition: I/O-intensive programs spend most of their time waiting for input/output operations to complete, during which time the CPU often sits idle.

Implementation:

```
void io_worker(void) {
    char filename[64], buffer[IO_FILE_SIZE];
    snprintf(filename, sizeof(filename), "/tmp/io_test_%d.tmp", getpid());
    for (int iter = 0; iter < BASE_LOOP_COUNT; iter++) {
        int fd = open(filename, O_WRONLY | O_CREAT | O_TRUNC, 0644);
        for(int k = 0; k < 10; k++) {
            write(fd, buffer, IO_FILE_SIZE);
            fsync(fd); // Force disk I/O
        }
        close(fd);
    }
    unlink(filename);
}
```

Line-by-Line Explanation:

- Unique filename per process: Uses `getpid()` or `pthread_self()` to avoid conflicts
- `open()` with `O_TRUNC`: Opens file and truncates it each iteration
- `write()`: Writes 1KB buffer to file
- `fsync()`: CRITICAL - Forces data to be written to physical disk, not just OS cache
- `unlink()`: Cleans up temporary file after completion

4. Part C: Measurement Script and Results

4.1 Shell Script Explanation

The bash script automates execution and metric collection for all six program variants.

Key Components:

CPU Pinning with taskset:

```
taskset -c 0 $bin $work
```

Pins the program to CPU 0 for consistent measurement. This ensures all processes/threads compete for the same CPU core during baseline measurement.

CPU and Memory Monitoring with top:

```
top -b -n 1 -p $p 2>/dev/null | tail -1
```

Runs top in batch mode (-b), takes 1 sample (-n 1), for specific PID (-p). The script aggregates CPU% and Mem% across all child processes/threads.

I/O Monitoring with iostat:

```
iostat -dx 1 > "$iostat_log" 2>&1 &
```

Runs iostat in extended mode (-x) for all devices (-d), sampling every 1 second, running in background.

Time Measurement:

```
/usr/bin/time -p -o "$time_log" $cmd &
```

Uses /usr/bin/time (not shell built-in) with POSIX format (-p), output to file (-o).

4.2 Part C Results

Program+Function	CPU%	Mem%	IO
A+cpu	100.00	0	15.905
B+cpu	102.23	0	15.905
A+mem	70.00	0.15	15.905
B+mem	66.66	0.13	15.905
A+io	1.91	0	4.54
B+io	1.21	0	3.98

4.3 Part C Analysis

CPU Workers (A+cpu, B+cpu):

- CPU% ≈ 100%: Both processes and threads fully utilize the CPU
- Mem% = 0: CPU worker only uses ~16 bytes (one double variable), which is negligible compared to system RAM

- This confirms the worker is truly CPU-bound

Memory Workers (A+mem, B+mem):

- CPU% \approx 66-70%: Lower than CPU workers because time is spent waiting for memory access
- Mem% \approx 0.13-0.15%: Each worker allocates \sim 40MB. On an 8GB system, this is \sim 0.5% per worker
- Random access pattern defeats CPU cache, causing memory bottleneck

I/O Workers (A+io, B+io):

- CPU% \approx 1-2%: Very low because the process spends most time waiting for disk
- Mem% = 0: Only uses 1KB buffer, negligible memory
- fsync() forces synchronous writes, making the disk the bottleneck

5. Part D: Scaling Analysis

5.1 Scaling Configuration

Part D extends Part C by varying the number of processes/threads:

- Program A (fork): 2, 3, 4, 5 processes
- Program B (pthread): 2, 3, 4, 5, 6, 7, 8 threads

Real-World CPU Pinning:

```
taskset -c 0-$((NUM_CPUS-1)) $bin $work $cnt
```

For scaling analysis, processes/threads are pinned to ALL available CPUs, allowing the OS scheduler to distribute work naturally. This represents real-world behaviour where applications don't manually pin threads to specific cores.

5.2 CPU Utilization Plot

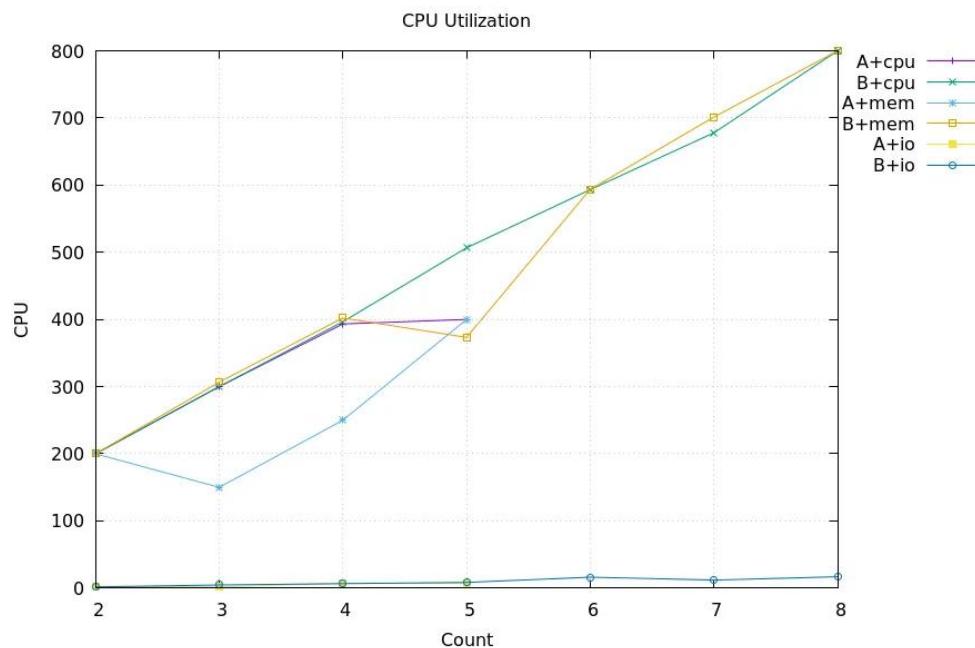


Figure 1: CPU Utilization vs Process/Thread Count

Analysis:

- CPU and Memory workers show linear scaling: 2 workers = ~200%, 4 workers = ~400%, 8 workers = ~800%
- I/O workers remain flat near 0% regardless of count (disk-bound, not CPU-bound)
- Threads (B) and Processes (A) scale similarly for CPU-bound work

5.3 Memory Utilization Plot

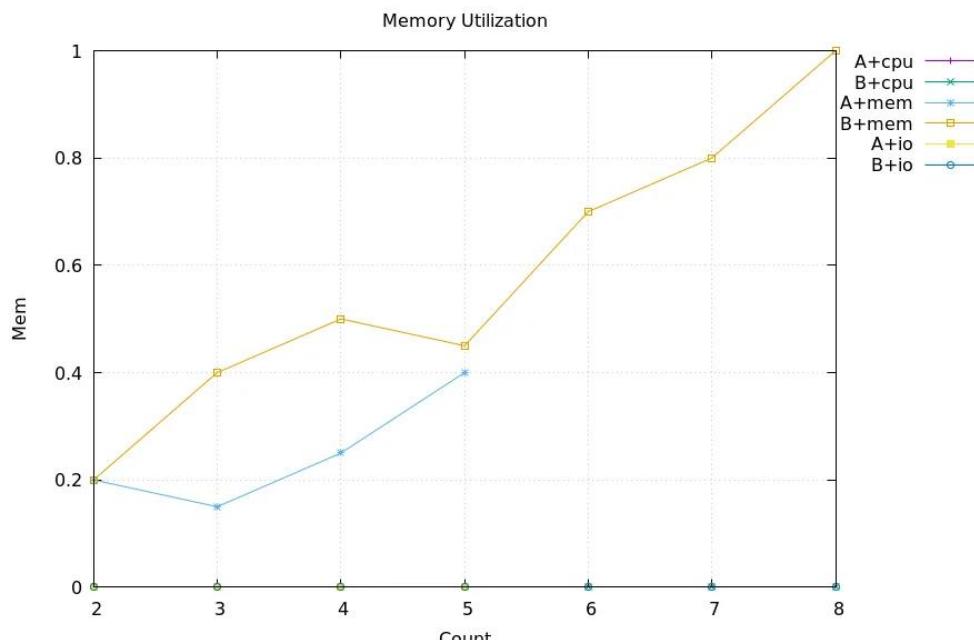


Figure 2: Memory Utilization vs Process/Thread Count

Analysis:

- Only memory workers (A+mem, B+mem) show significant memory usage
- B+mem scales from 0.2% to 1.0% as threads increase from 2 to 8
- CPU and I/O workers show 0% memory (they don't allocate significant RAM)
- Threads share address space but each allocates its own 40MB array

5.4 I/O Activity Plot

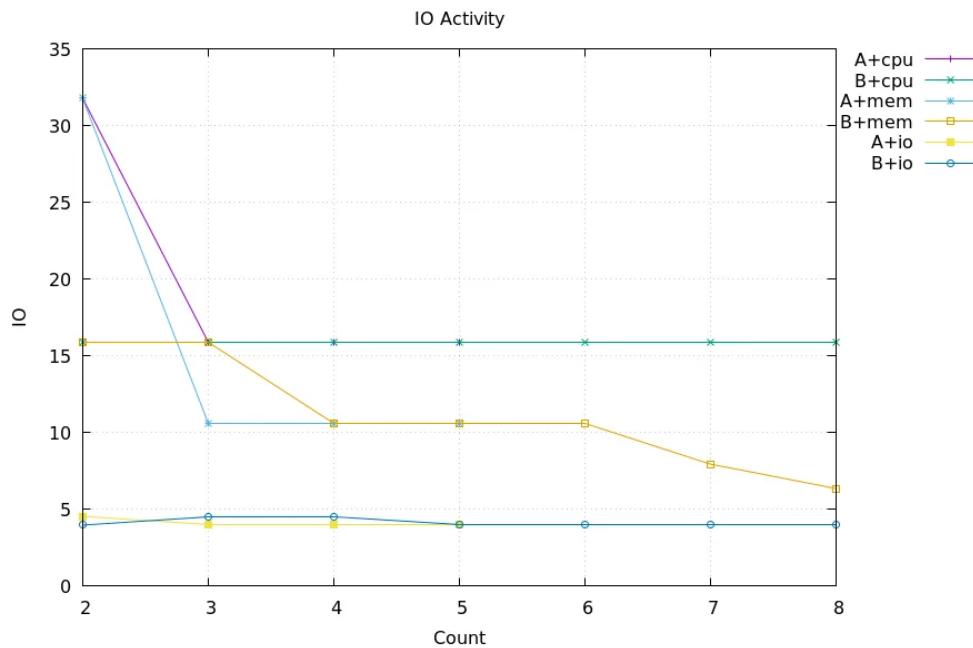


Figure 3: I/O Activity vs Process/Thread Count

Analysis:

- I/O activity is measured as average KB/s (read + write)
- I/O workers show relatively flat I/O because disk bandwidth is saturated
- Adding more workers doesn't increase total I/O throughput (disk bottleneck)
- The ~15 KB/s baseline for CPU/mem workers represents system background I/O

5.5 Execution Time Plot

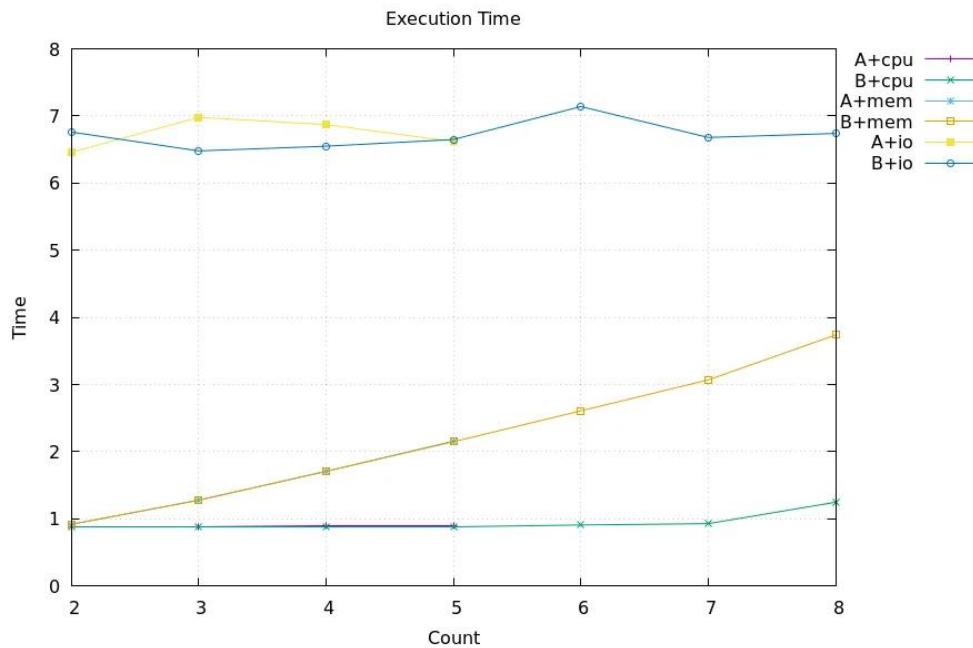


Figure 4: Execution Time vs Process/Thread Count

Analysis:

- CPU workers (A+cpu, B+cpu): Near-constant time (~0.9s) - true parallel execution
- Memory workers (B+mem): Time increases with count due to cache contention and memory bandwidth limits
- I/O workers: Constant high time (~6.5s) regardless of count - disk is the bottleneck
- This demonstrates that adding parallelism doesn't help I/O-bound workloads

6. Observations and Analysis

6.1 Process vs Thread Comparison

Aspect	Processes (fork)	Threads (pthread)
Address Space	Separate (copied)	Shared
Creation Overhead	Higher (copy pages)	Lower (shared pages)
Communication	IPC required	Direct memory access
Fault Isolation	Better (separate space)	Worse (shared space)

6.2 Key Findings

1. CPU-Bound Workloads:

Both processes and threads achieve near-perfect CPU utilization. The CPU% scales linearly with count until hardware limits are reached. Execution time remains constant with parallel execution.

2. Memory-Bound Workloads:

Memory workers show lower CPU% due to memory access latency. Random access patterns defeat CPU caches. Time increases with more workers due to memory bandwidth contention.

3. I/O-Bound Workloads:

I/O workers show very low CPU% (1-2%) as they spend most time waiting. Adding more workers doesn't reduce execution time. The disk (not CPU) is the bottleneck. fsync() ensures synchronous writes, making this truly I/O-bound.

4. Scaling Behavior:

CPU-bound: Linear scaling up to number of cores. Memory-bound: Sub-linear scaling due to shared memory bus. I/O-bound: No scaling benefit (single disk bottleneck).

7. AI Usage Declaration

AI USAGE DECLARATION

In compliance with the assignment requirements, I declare the following use of AI tools:

7.1 AI Tool Used

Tool: Claude (Anthropic)

Purpose: Code assistance, debugging, and Plot analysis

7.2 AI-Assisted Components

1. Code Structure and Implementation:

- Initial structure of Program A (fork-based) and Program B (pthread-based)
- Worker function implementations (cpu, mem, io)
- Shell scripts for automation (Part C and Part D)

2. Debugging and Optimization:

- Fixing process ID tracking in shell scripts (finding real PID vs /usr/bin/time PID)
- Optimizing loop counts for visibility in 'top'
- Discussion on CPU pinning strategies (real-world vs controlled)

3. Report Generation:

- Plot analysis and explanations in the report were developed through discussion with AI

7.3 Human Contributions

- Understanding and verification of all code
- Running experiments and collecting actual measurements
- Decision-making on implementation approaches
- Verification of results and analysis

7.4 Verification Statement

I confirm that I understand every line of code submitted and can explain its functionality during the viva examination.

8. Conclusion

This assignment provided hands-on experience with process and thread-based parallelism in Linux. The key takeaways are:

1. Processes (fork) and threads (pthread) offer different trade-offs in terms of isolation vs. efficiency
2. CPU-bound workloads benefit most from parallelism, achieving near-linear scaling
3. Memory-bound workloads show diminishing returns due to shared memory bandwidth
4. I/O-bound workloads cannot be sped up through CPU parallelism alone
5. Proper measurement requires understanding of tools like top, taskset, iostat, and time

GitHub Repository: https://github.com/WiseShukla/GRS_Assignment1

9. Appendix: Complete Source Code

9.1 Makefile

```
CC = gcc
CFLAGS = -Wall -Wextra -O2
LDFLAGS_A = -lm
LDFLAGS_B = -lpthread -lm

all: program_a program_b

program_a: PhD25001_Part_A_Program_A.c
    $(CC) $(CFLAGS) -o program_a PhD25001_Part_A_Program_A.c $(LDFLAGS_A)

program_b: PhD25001_Part_B_Program_B.c
    $(CC) $(CFLAGS) -o program_b PhD25001_Part_B_Program_B.c $(LDFLAGS_B)

clean:
    rm -f program_a program_b *.csv *.png *.tmp
```

Note: Complete source code files (Program A, Program B, Shell Scripts) are included in the submission folder.