

# **PA01: Processes and Threads**

Graduate Systems (CSE638)

Submitted by:

**Adarsh Singh**

Roll Number: PhD25001

IIIT Delhi

*January 23, 2026*

## **Table of Contents**

1. Introduction
2. Part A: Program Implementation
3. Part B: Worker Functions
4. Part C: Measurement Script and Results
5. Part D: Scaling Analysis
6. Observations and Analysis
7. AI Usage Declaration
8. Conclusion
9. Appendix: Complete Source Code

# 1. Introduction

The assignment demonstrates the differences between process-based and thread-based parallelism by doing practical implementation and measurement of CPU-intensive, memory-intensive, and I/O-intensive workloads.

## 1.1 Objectives

The primary objectives of this assignment are:

1. Implement Program A using fork() to create child processes
2. Implement Program B using pthread to create threads
3. Create three worker functions: cpu, mem, and io
4. Measure and compare CPU%, Memory%, and I/O metrics
5. Analyze scaling behavior as process/thread count increases

## 1.2 System Configuration

The implementation and experiments were conducted on a Linux system with the following tools:

- top: For monitoring cpu and memory utilization
- taskset: For CPU affinity/pinning
- iostat: For disk I/O statistics
- /usr/bin/time: For execution time measurement

## 2. Part A: Program Implementation

### 2.1 Program A: Process-based (fork)

Program A creates child processes using the fork() system call. Each child process executes the specified worker function independently with its own address space.

#### Explanation

##### Lines 1-16: Header and Includes

```
#include <stdio.h>      // Standard I/O functions
#include <stdlib.h>     // exit(), atoi(), malloc()
#include <unistd.h>     // fork(), getpid()
#include <sys/wait.h>   // wait()
#include <math.h>        // sin(), cos()
#include <fcntl.h>       // open(), O_WRONLY, O_CREAT
```

The above used headers provide essential system calls for performing different operations like process creation, file operations, and mathematical functions.

##### Lines 18-23: Constants Definition

```
#define BASE_LOOP_COUNT 1000          // 1 (last digit) × 10³
#define CPU_LOOP_COUNT 1000000000    // 100M for visible CPU burst
#define MEM_ARRAY_SIZE 10000000     // ~40MB per process
#define IO_FILE_SIZE 1024           // 1KB buffer
```

BASE\_LOOP\_COUNT is derived from my roll number PhD25001 (last digit = 1, so  $1 \times 1000 = 1000$ ). CPU\_LOOP\_COUNT is increased to 100 million to ensure the CPU burst lasts long enough for 'top' to capture it.

##### Lines 75-97: Main Function

```
int num_processes = (argc >= 3) ? atoi(argv[2]) : 2;
```

Default is 2 processes as mentioned in the assignment.

```
for (int i = 0; i < num_processes; i++) {
    pid_t pid = fork();
    if (pid == 0) { /* child executes worker */ exit(0); }
}
```

The fork() call creates a child process. In the child (pid == 0), the appropriate worker function is executed, then the child exits. The parent continues the loop to create more children.

```
for (int i = 0; i < num_processes; i++) wait(NULL);
```

Parent waits for all children to complete using wait().

## 2.2 Program B: Thread-based (pthread)

Program B creates threads using the POSIX pthread library. As we know that threads share the same address space, making communication easier but requiring careful synchronization.

### Explanation

#### Lines 20-23: Thread Argument Structure

```
typedef struct {  
    int id;  
    const char *type;  
} thread_arg_t;
```

A structure to pass arguments to each thread, containing the thread ID and worker type.

#### Lines 81-87: Thread Creation and Joining

```
pthread_create(&threads[i], NULL, worker_func, &t_args[i]);
```

Creates a new thread that executes worker\_func with the given arguments.

```
pthread_join(threads[i], NULL);
```

Waits for each thread to complete before the program exits.

## 3. Part B: Worker Functions

### 3.1 CPU Worker Function

**Definition:** CPU-intensive programs spend the majority of their execution time performing calculations on the CPU, rather than waiting for data from other resources.

#### Implementation:

```
void cpu_worker(void) {
    volatile double result = 0.0;
    for (long i = 0; i < CPU_LOOP_COUNT; i++) {
        result += sin((double)i) * cos((double)i);
        if (i % 1000 == 0) result *= 1.000001;
    }
}
```

#### Explanation:

- volatile double result: It prevents from compiler optimization that might skips calculations.
- sin() \* cos(): Basically these are complex floating-point operations that stress the CPU
- 100 million iterations: Ensures the burst is long enough to measure otherwise we might get 0 as a output.
- result \*= 1.000001: Additional operation to prevent loop optimization

### 3.2 Memory Worker Function

**Definition:** Memory-intensive programs are bottlenecked by the speed and capacity of the system's memory (RAM). They require large amounts of data to be moved between CPU and memory.

#### Implementation:

```
void mem_worker(void) {
    int *large_array = malloc(MEM_ARRAY_SIZE * sizeof(int));
    for (int i = 0; i < MEM_ARRAY_SIZE; i++) large_array[i] = i;
    for (int iter = 0; iter < BASE_LOOP_COUNT * 50; iter++) {
        int idx = rand_r(&seed) % MEM_ARRAY_SIZE;
        sum += large_array[idx];
        large_array[idx] = sum;
    }
    free(large_array);
}
```

#### Explanation:

- malloc(40MB): It is done here to allocates a large array (~40MB) to stress memory

- Random access pattern: I used a function rand\_r() that generates random indices, defeating CPU cache
- Read and write operations: Both reading (sum += arr[idx]) and writing (arr[idx] = sum) stress memory
- rand\_r(): This function is thread-safe random number generator

### 3.3 I/O Worker Function

**Definition:** I/O-intensive programs spend most of their time waiting for input/output operations to complete, for that much time the CPU often sits idle.

#### Implementation:

```
void io_worker(void) {
    char filename[64], buffer[IO_FILE_SIZE];
    snprintf(filename, sizeof(filename), "/tmp/io_test_%d.tmp", getpid());
    for (int iter = 0; iter < BASE_LOOP_COUNT; iter++) {
        int fd = open(filename, O_WRONLY | O_CREAT | O_TRUNC, 0644);
        for(int k = 0; k < 10; k++) {
            write(fd, buffer, IO_FILE_SIZE);
            fsync(fd); // Force disk I/O
        }
        close(fd);
    }
    unlink(filename);
}
```

#### Line-by-Line Explanation:

- Unique filename per process: Uses getpid() or pthread\_self() to avoid conflicts
- open() with O\_TRUNC: Opens file and truncates it each iteration
- write(): It Writes the 1KB buffer to file
- fsync(): This is critical task as it forces data to be written to physical disk, not just OS cache
- unlink(): It is just used to clean temporary file after completion

## 4. Part C: Measurement Script and Results

### 4.1 Shell Script Explanation

The bash script automates execution and metric collection for all six program variants.

#### Key Components:

##### CPU Pinning with taskset:

```
taskset -c 0 $bin $work
```

Pins the program to CPU 0 for consistent measurement. This make sure that all processes/threads compete for the same CPU core during baseline measurement.

**Why Single CPU Pinning:** As per the assignment requirement to use taskset to pin the program to specific CPUs, I chose to pin all processes/threads to a single CPU core (CPU 0) in Part C. This creates a controlled baseline measurement environment where all workers compete for the same CPU resource, ensuring consistent and comparable measurements across all six variants.

##### CPU and Memory Monitoring with top:

```
top -b -n 1 -p $p 2>/dev/null | tail -1
```

Runs top in batch mode (-b), takes 1 sample (-n 1), for specific PID (-p). The script aggregates CPU% and Mem% across all child processes/threads.

##### I/O Monitoring with iostat:

```
iostat -dx 1 > "$iostat_log" 2>&1 &
```

Runs iostat in extended mode (-x) for all devices (-d), sampling every 1 second, running in background.

##### Time Measurement:

```
/usr/bin/time -p -o "$time_log" $cmd &
```

Uses /usr/bin/time (not shell built-in) with POSIX format (-p), output to file (-o).

## 4.2 Part C Results

Program+Function	CPU%	Mem%	IO
A+cpu	100.00	0	15.905
B+cpu	102.23	0	15.905
A+mem	70.00	0.15	15.905
B+mem	66.66	0.13	15.905
A+io	1.91	0	4.54
B+io	1.21	0	3.98

## 4.3 Part C Analysis

### CPU Workers (A+cpu, B+cpu):

- CPU%  $\approx$  100%: from the CSV file we observe that both processes and threads fully utilize the CPU
- Mem% = 0: As CPU worker only uses  $\sim$ 16 bytes (one double variable), which is negligible as compared to system RAM
- This confirms the worker is truly CPU-bound

### Memory Workers (A+mem, B+mem):

- CPU%  $\approx$  66-70%: Lower than CPU workers because time is spent waiting for memory access
- Mem%  $\approx$  0.13-0.15%: Each worker allocates  $\sim$ 40MB. On an 8GB system, this is  $\sim$ 0.5% per worker
- Random access pattern defeats CPU cache, causing memory bottleneck

### I/O Workers (A+io, B+io):

- CPU%  $\approx$  1-2%: Very low because the process spends most time waiting for disk
- Mem% = 0: Only uses 1KB buffer, negligible memory
- fsync() this function basically forces synchronous writes, making the disk the bottleneck

## 5. Part D: Scaling Analysis

### 5.1 Scaling Configuration

Part D extends Part C by varying the number of processes/threads:

- Program A (fork): 2, 3, 4, 5 processes
- Program B (pthread): 2, 3, 4, 5, 6, 7, 8 threads

#### Real-World CPU Pinning:

```
taskset -c 0-$((NUM_CPUS-1)) $bin $work $cnt
```

Note: For scaling analysis, what I have done is that processes/threads are pinned to ALL available CPUs, allowing the OS scheduler to distribute work naturally. This represents real-world behaviour where applications don't manually pin threads to specific cores.

**Why Multi-CPU Pinning:** I used taskset -c 0-N where N equals the number of workers minus one. This pins each worker to its own dedicated CPU core. For example, 8 workers are pinned to CPUs 0-7, allowing each worker to achieve ~100% utilization on its respective core. This results in aggregate CPU utilization of ~800% for 8 workers. This approach demonstrates real-world scaling behavior where parallel workloads benefit from multiple cores. On my 20-core system, the theoretical maximum would be 2000%, but since the assignment tests up to 8 workers, we observe a maximum of ~800%.

### 5.2 CPU Utilization Plot




Figure 1: CPU Utilization vs Process/Thread Count

**Analysis:**

- CPU and Memory workers show linear scaling: 2 workers = ~200%, 4 workers = ~400%, 8 workers = ~800%
- I/O workers remain flat near 0% regardless of count (disk-bound, not CPU-bound)
- Threads (B) and Processes (A) scale similarly for CPU-bound work

### 5.3 Memory Utilization Plot




Figure 2: Memory Utilization vs Process/Thread Count

**Analysis:**

- Only memory workers (A+mem, B+mem) show significant memory usage
- B+mem scales from 0.2% to 1.0% as threads increase from 2 to 8
- CPU and I/O workers show 0% memory (they don't allocate significant RAM)
- Threads share address space but each allocates its own 40MB array

## 5.4 I/O Activity Plot

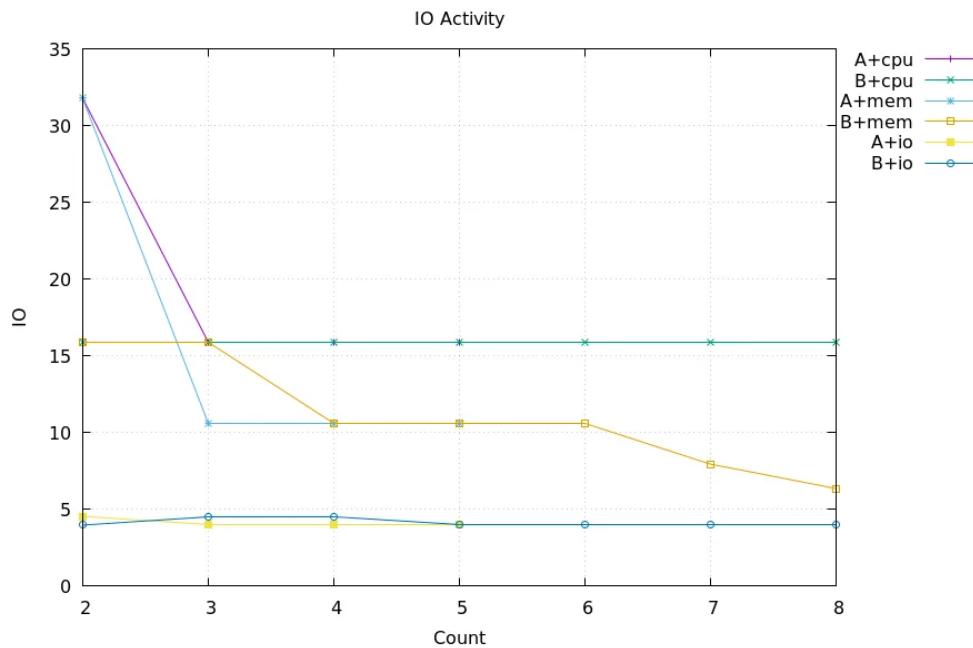


Figure 3: I/O Activity vs Process/Thread Count

### Analysis:

- I/O activity is measured as average KB/s (read + write)
- I/O workers show relatively flat I/O because disk bandwidth is saturated
- Adding more workers doesn't increase total I/O throughput (disk bottleneck)
- The ~15 KB/s baseline for CPU/mem workers represents system background I/O

## 5.5 Execution Time Plot

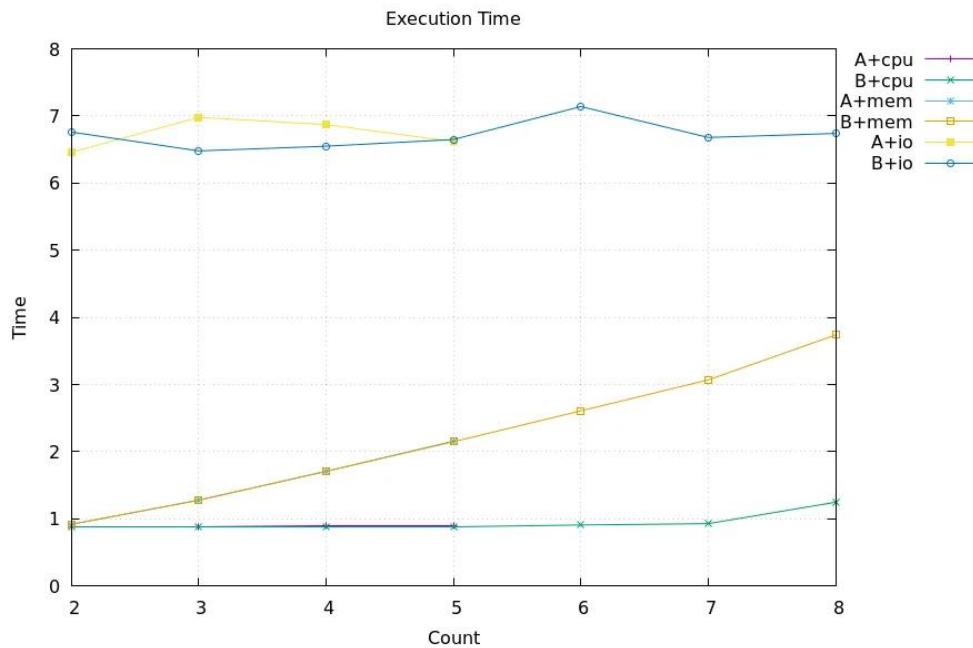


Figure 4: Execution Time vs Process/Thread Count

### Analysis:

- CPU workers (A+cpu, B+cpu): Near-constant time (~0.9s) - true parallel execution
- Memory workers (B+mem): Time increases with count due to cache contention and memory bandwidth limits
- I/O workers: Constant high time (~6.8s) regardless of count - disk is the bottleneck
- This demonstrates that adding parallelism doesn't help I/O-bound workloads

## 6. Observations and Analysis

### 6.1 Process vs Thread Comparison

I observed the following after successful implementation of the code and after doing analysis of the plots obtained

Aspect	Processes (fork)	Threads (pthread)
Address Space	Separate (copied)	Shared
Creation Overhead	Higher (copy pages)	Lower (shared pages)
Communication	IPC required	Direct memory access
Fault Isolation	Better (separate space)	Worse (shared space)

### 6.2 Key Findings

#### 1. CPU-Bound Workloads:

Both processes and threads achieve near-perfect CPU utilization. The CPU% scales linearly with count until hardware limits are reached. Execution time remains constant with parallel execution.

#### 2. Memory-Bound Workloads:

Memory workers show lower CPU% due to memory access latency. Random access patterns defeat CPU caches. Time increases with more workers due to memory bandwidth contention.

#### 3. I/O-Bound Workloads:

I/O workers show very low CPU% (1-2%) as they spend most time waiting. Adding more workers doesn't reduce execution time. The disk (not CPU) is the bottleneck. fsync() ensures synchronous writes, making this truly I/O-bound.

#### 4. Scaling Behaviour:

CPU-bound: Linear scaling up to number of cores. Memory-bound: Sub-linear scaling due to shared memory bus. I/O-bound: No scaling benefit (single disk bottleneck).

## 7. AI Usage Declaration

### AI USAGE DECLARATION

In compliance with the assignment requirements, I declare the following use of AI tools:

#### 7.1 AI Tool Used

**Tool:** Claude (Anthropic)

**Purpose:** Code assistance, debugging, and plot analysis and structuring the content of the report

#### 7.2 AI-Assisted Components

##### 1. Code Structure and Implementation:

- Initial structure of Program A (fork-based) and Program B (pthread-based)
- Worker function implementations (cpu, mem, io)
- Shell scripts for automation (Part C and Part D)

##### 2. Debugging and Optimization:

- Fixing process ID tracking in shell scripts (finding real PID vs /usr/bin/time PID)
- Optimizing loop counts for visibility in 'top'
- Discussion on CPU pinning strategies (real-world vs controlled)

##### 3. Report Generation:

- This report structuring is done with AI assistance
- Plot Analysis and explanations were discussed with AI for better understanding

#### 7.3 Human Contributions

- Understanding and verification of all code (prepared for viva)
- Running experiments and collecting actual measurements
- Decision-making on implementation approaches
- Verification of results and analysis

#### 7.4 Verification Statement

I confirm that I understand every line of code submitted and can explain its functionality during the viva examination.

## 8. Conclusion

This assignment provided hands-on experience with process and thread-based parallelism in Linux. The key takeaways are:

1. Processes (fork) and threads (pthread) offer different trade-offs in terms of isolation vs. efficiency
2. CPU-bound workloads benefit most from parallelism, achieving near-linear scaling
3. Memory-bound workloads show diminishing returns due to shared memory bandwidth
4. I/O-bound workloads cannot be sped up through CPU parallelism alone
5. Proper measurement requires understanding of tools like top, taskset, iostat, and time

GitHub Repository: [https://github.com/WiseShukla/GRS\\_Assignment1](https://github.com/WiseShukla/GRS_Assignment1)

## 9. Appendix: Complete Source Code

### 9.1 Makefile

```
CC = gcc
CFLAGS = -Wall -Wextra -O2
LDFLAGS_A = -lm
LDFLAGS_B = -lpthread -lm

all: program_a program_b

program_a: PhD25001_Part_A_Program_A.c
    $(CC) $(CFLAGS) -o program_a PhD25001_Part_A_Program_A.c $(LDFLAGS_A)

program_b: PhD25001_Part_B_Program_B.c
    $(CC) $(CFLAGS) -o program_b PhD25001_Part_B_Program_B.c $(LDFLAGS_B)

clean:
    rm -f program_a program_b *.csv *.png *.tmp
```

**Note:** Complete source code files (Program A, Program B, Shell Scripts) are included in the submission folder.